

Efficient green energy sources report

Alejandro Barajas, Gary Layman, Trevor Olmo, Add your names here

2025-10-14

R and Python library set up

- **Overall goal** The goal is to identify which sources deliver the most energy per dollar while minimizing environmental impact

Data Preparation

```
#### Commented this out because we no longer need it
# energy <- eia_data(
#   dir = "electricity/electric-power-operational-data",
#   data = c("ash-content", "consumption-for-eg", "consumption-for-eg-btu", "consumption-uto"),
#   freq = "quarterly",
#   start = "2020",
#   end = "2025"
#)
# All columns are <char> and will need to be cleaned

# YOUR CODE HERE
```

Data Collection & Cleaning

- What sources are used?
- What cleaning steps are necessary?

We already scraped the data from the EIA website, so we can use that data directly. We just need to read it in from the csvs it has been saved in.

```
# YOUR CODE HERE
#Find all csv files in the directory, read them in, and store them in a single df
eia_csvs = glob("datasets/*.csv")
dfs = [pd.read_csv(csv) for csv in eia_csvs]
energy_df = pd.concat(dfs, ignore_index=True)
energy_df.head()
```

	period	location	...	total-consumption-btu	total-consumption-btu-units
0	2016-Q3	IN	...	8.93140	million MMBtu
1	2016-Q3	IN	...	0.31604	million MMBtu
2	2016-Q3	MD	...	NaN	million MMBtu
3	2016-Q3	MD	...	0.00000	million MMBtu
4	2016-Q3	MD	...	1.65269	million MMBtu

[5 rows x 37 columns]

Data Organization & Conditioning

- How is the data structured?
- Are there transformations or feature engineering steps?

Here I just check number of rows and columns.

```
# YOUR CODE HERE
# Number of rows
print("Rows: ", energy_df.shape[0])
# Number of columns
print("Columns: ", energy_df.shape[1])
```

Rows: 530986

Columns: 37

Columns Types

```
#Data types
energy_df.dtypes
```

period	object
location	object
stateDescription	object

sectorid	int64
sectorDescription	object
fueltypeid	object
fuelTypeDescription	object
ash-content	float64
ash-content-units	object
consumption-for-eg	float64
consumption-for-eg-units	object
consumption-for-eg-btu	float64
consumption-for-eg-btu-units	object
consumption-uto	float64
consumption-uto-units	object
consumption-uto-btu	float64
consumption-uto-btu-units	object
cost	float64
cost-units	object
cost-per-btu	float64
cost-per-btu-units	object
generation	float64
generation-units	object
heat-content	float64
heat-content-units	object
receipts	float64
receipts-units	object
receipts-btu	float64
receipts-btu-units	object
stocks	float64
stocks-units	object
sulfur-content	float64
sulfur-content-units	object
total-consumption	float64
total-consumption-units	object
total-consumption-btu	float64
total-consumption-btu-units	object
dtype:	object

Data Storage

- Where and how is the data stored?
- Is it accessible and secure?

Find the number of missing values in each column.

```
# YOUR CODE HERE
energy_df.isna().sum()
```

```
period                0
location              0
stateDescription      0
sectorid              0
sectorDescription     0
fueltypeid            0
fuelTypeDescription   0
ash-content           189444
ash-content-units     0
consumption-for-eg    49532
consumption-for-eg-units 0
consumption-for-eg-btu 49532
consumption-for-eg-btu-units 0
consumption-uto       49532
consumption-uto-units 0
consumption-uto-btu   49532
consumption-uto-btu-units 0
cost                  433700
cost-units            0
cost-per-btu          485548
cost-per-btu-units    0
generation            15880
generation-units      0
heat-content          170783
heat-content-units    0
receipts              170783
receipts-units        0
receipts-btu          166205
receipts-btu-units    0
stocks                507804
stocks-units          0
sulfur-content        187923
sulfur-content-units  0
total-consumption     49532
total-consumption-units 0
total-consumption-btu 49532
total-consumption-btu-units 0
dtype: int64
```

Shows the fuel types that have missing cost values

```
# YOUR CODE HERE
```

```
missing_cost = energy_df[energy_df["cost"].isna()]
missing_cost["fuelTypeDescription"].unique()
```

```
array(['renewable', 'petroleum', 'residual fuel oil',
      'natural gas & other gases', 'natural gas', 'fossil fuels',
      'distillate fuel oil', 'all coal products',
      'coal, excluding waste coal', 'bituminous coal',
      'waste oil and other oils', 'onshore wind turbine', 'biomass',
      'other renewables', 'wood and wood wastes', 'wind',
      'renewable waste products', 'bituminous coal and synthetic coal',
      'all renewables', 'all fuels', 'solar', 'subbituminous coal',
      'solar photovoltaic', 'refined coal', 'petroleum liquids', 'other',
      'biogenic municipal solid waste', 'municipal landfill gas',
      'landfill gas', 'conventional hydroelectric', 'petroleum coke',
      'nuclear', 'lignite coal', 'hydro-electric pumped storage',
      'solar thermal', 'waste coal', 'other gases', 'geothermal',
      'anthracite coal', 'offshore wind turbine'], dtype=object)
```

Finds the counts for each fuel type

```
# YOUR CODE HERE
```

```
missing_cost["fuelTypeDescription"].value_counts()
```

fuelTypeDescription	
biomass	38855
all fuels	23127
fossil fuels	21748
natural gas & other gases	20655
renewable	20351
all renewables	19817
petroleum	19020
distillate fuel oil	18418
renewable waste products	15839
natural gas	15181
petroleum liquids	13586
coal, excluding waste coal	12871
solar photovoltaic	12593
solar	12580

other	12218
municipal landfill gas	11917
landfill gas	11126
conventional hydroelectric	11015
bituminous coal	10502
bituminous coal and synthetic coal	10489
wind	10406
onshore wind turbine	10318
other renewables	10218
wood and wood wastes	9830
all coal products	9372
residual fuel oil	8522
waste oil and other oils	8066
subbituminous coal	7656
nuclear	4644
biogenic municipal solid waste	4461
other gases	4277
refined coal	3267
hydro-electric pumped storage	2614
waste coal	2071
petroleum coke	1986
geothermal	1487
lignite coal	1338
solar thermal	796
offshore wind turbine	236
anthracite coal	227

Name: count, dtype: int64

Exploratory Data Analysis (EDA)

- What are the high-level patterns?
- What intuitive insights emerge?
- Is the data consistent across sources?
- What assumptions are we making?
- What new questions arise from the EDA?

Note: Be mindful of *confirmation bias*—avoid interpreting data only to support your initial hypothesis.

```
# YOUR CODE HERE
```

```
# YOUR CODE HERE
```

Model Planning

- **Model Type:** Classification, clustering, regression—what fits best?
- **Success/Failure Definitions:** Refine what constitutes a good or bad outcome.
- **EDA for Modeling:**
 - Identify relevant variables
 - Explore correlations
 - Apply domain knowledge
 - Verify assumptions
 - Prototype modeling ideas

```
# YOUR CODE HERE
```

```
# YOUR CODE HERE
```

Model Building

- What models are being trained?
- How are they tested?
- What metrics are used to evaluate performance?

```
# YOUR CODE HERE
```

```
# YOUR CODE HERE
```

Communicating Results

- Are the results **robust** across different scenarios?
- Are they **statistically significant**?
- Why do these results **matter** to stakeholders?
- How do they compare to your definitions of success and failure?

Operationalization

- **Presentation:**
 - Who is the audience?
 - Is the code and analysis well-documented and replicable?
- **Deployment:**
 - Are models deployed on live data?
 - Are they behaving as expected?